

Capstone Project - Car Accident Severity (Week 3)

Python Notebook

Table of contents

- [Introduction: Business Problem](#)
- [Data Preparation](#)
- [Methodology](#)
- [Exploratory Data Analysis](#)
- [Modeling, Testing and Evaluation](#)
- [Results and Discussion](#)
- [Conclusion and Outlook](#)

Introduction: Business Problem

1. Introduction / Business Problem

1.1. Background

Road traffic injuries (RTIs) are a major public health problem. The annual global status reports on road safety, launched by the World Health Organization (WHO), highlights that the number of road traffic deaths has exceeded one million in recent years. That is over 3000 people dying on the world's roads every day.[1] Therefore, analyzing the various factors that could help predict accident severity can guide the government administration to implement changes in a timely manner that may reduce the number of fatalities & serious injuries.

In the past few years, the volume of research in the areas of accident analysis and prediction has been increasing. Among the analytical data mining solutions, supervised machine learning (ML), has become a popular scientific method to predict the severity of accidents. The reasons for this popularity, can be referred to the capacity present in ML to identify the existing patterns in the data and make predictions via the establishment and evaluation of diverse algorithms. Moreover, the ability of MLs to handle large amounts of data is an additional asset for this purpose, as the data on road traffic accidents are often sparse and largely extended.

1.2. Objective

The objective of this capstone project is to analyze the collision data set for Seattle, WA and determine the most possible factors including weather, road conditions, visibility, and various other factors that best predict accident severity by training and evaluating supervised machine learning algorithms.

This project will be used to answer the business question: How can the city of Seattle, Washington best predict the severity of collisions that occur and what avenues can be explored to remedy this issue?

1.3. Target Audience

The report of this project can be targeted to stakeholders, who are involved with road traffic injuries, such as road administrators, traffic control authorities, and emergency road services in order to help them predict the car accident severities and improve the road users' safety margins.

References

[1].World Health Organization: <http://www.who.int>

Data Preparation

2.1. Data Source

It is now time to understand the data and then prepare it to be fed into the modeling tools. The given dataset used in this project (provided by the coursera example data) can be downloaded

here: <https://s3.us.cloud-object-storage.appdomain.cloud/cf-courses-data/CognitiveClass/DP0701EN/version-2/Data-Collisions.csv>

The example dataset (Data-Collisions.csv) contains 194673 and 38 columns including the labeled data. The labeled data is the "Severity Code", which describes the fatality of an accident. In the shared dataset, the severity code column consists of two values: 1 for property damage and 2 for injury. The dataset includes different attributes, describing a variety of conditions: location, weather, light, road, collision types, and so forth that may influence the severity of the accidents. The attributes are of the types of int64, float64, or object.

Lets first load required libraries:

In [86]:

```
import pandas as pd
import numpy as np
from sklearn import metrics
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
from matplotlib.ticker import NullFormatter
import matplotlib.ticker as ticker
from sklearn import preprocessing
%matplotlib inline
from sklearn.utils import resample
from sklearn.ensemble import ExtraTreesClassifier
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.feature_selection import SelectFromModel
from sklearn.metrics import accuracy_score
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_curve, auc
import matplotlib.image as mpimg
from sklearn import tree
from sklearn.tree import export_graphviz
from sklearn import svm
from sklearn.ensemble import RandomForestClassifier
import matplotlib as mpl
```

Load Data From CSV File

In [2]:

```
df = pd.read_csv('Data-Collisions.csv')
df.head()
```

```
/Users/ZhouHui/opt/anaconda3/lib/python3.8/site-packages/IPython/core/interactiveshell.py:3071: DtypeWarning: Columns (33) have mixed types.Specify dt
ype option on import or set low_memory=False.
```

```
has_raised = await self.run_ast_nodes(code_ast.body, cell_name,
```

Out[2]:

[illegible]

[illegible]

[illegible]

distribution ratio is considerably high, the replacement of the missing data with reasonable new values is a better option as far as possible.

The other issue is the presence of both numerical and categorical data in the dataset. To this effect, the replacement is done by the frequency for the categorical variables and by mean for the numerical values. The missing categorical values that are replaced by the largest frequency belongs to the columns WEATHER, SPEEDING, LIGHTCOND, ROADCOND, JUNCTIONTYPE, INATTENTIONIND, COLLISIONTYPE, and ADDRTYPE. The missing numeric values in columns X and Y are replaced by the mean of the belonging columns, respectively.

What should be also taken into account, specifically for processing the data in the next steps, is the incompatibility of categorical variables with the predictive model analysis tools. For example, to develop regression models and being able to use packages such as Sklearn, these variables are converted into indicator variables during the data cleaning after handling the missing data.

In [7]:

```
# Rename X and Y with Longitude and Latitude
df1 = df.rename(columns={'X': 'LONGITUDE', 'Y': 'LATITUDE'})
df1.head()
```

Out [7] :

[illegible]

[illegible]

[illegible]

[illegible]

renamed to LONGITUDE and LATITUDE. The same analogy is also considered for the columns STATUS, INCDATE, INCDTTM, SDOT_COLDESC, PEDROWNOTGRNT, ST_COLDESC, PEDCYLCOUNT, HITPARKEDCAR, SEVERITYDESC, ADDRTYPE, and PEDCOUNT. The 10 features selected at the end of this step are listed in Table 1.

Table 1 - List of features being selected in the feature selection

| ID | Feature | Description |
|-----------|----------------|--|
| 1 | LONGITUDE | longitude |
| 2 | LATITUDE | latitude |
| 3 | PERSONCOUNT | total number of people involved in the collision |
| 4 | VEHCOUNT | the number of vehicles involved in the collision |
| 5 | JUNCTIONTYPE | category of junction at which collision took place |
| 6 | INATTENTIONIND | whether or not collision was due to |

| ID | Feature | Description |
|-----------|----------------|--|
| | | inattention |
| 7 | WEATHER | a description of the weather conditions during the time of the collision |
| 8 | ROADCOND | the condition of the road during the collision |
| 9 | LIGHTCOND | the light conditions during the collision. |
| 10 | SPEEDING | whether or not speeding was a factor in the collision |

In [8]:

```
# Drop LOCATION; Longitude and Latitude used instead.
# Two copies of SEVERITYCODE exist, drop the second SEVERITYCODE.1
#Drop columns includng codes: OBJECTID, INCKEY, COLDETKEY, REPORTNO,INTKEY
,EXCEPTRSNCODE, SDOT_COLCODE, SDOTCOLNUM --->
#ST_COLCODE, ST_COLDESC, SEGLANEKEY, CROSSWALKKEY
```

```
#Drop redundant info: STATUS, EXCEPTRSNDESC, INCDATE , INCDTTM, SDOT_COLDESC,
PEDROWNOTGRNT, ST_COLDESC, UNDERINFL--->
#PEDCYLCOUNT, HITPARKEDCAR, SEVERITYDESC, ADDRTYPE
df2 = df1.drop(["LOCATION", "SEVERITYCODE.1", "OBJECTID", "INCKEY", "COLDETKEY", "REPORTNO", "INTKEY",
               "EXCEPTRSNCODE", "SDOT_COLCODE", "ST_COLCODE", "SEGLANEKEY", "CROSSWALKKEY", "SDOTCOLNUM",
               "STATUS", "EXCEPTRSNDESC", "INCDATE", "INCDTTM", "SDOT_COLDESC", "PEDROWNOTGRNT", "UNDERINFL",
               "PEDCYLCOUNT", "HITPARKEDCAR", "ST_COLDESC", "SEVERITYDESC", "ADDRTYPE", "COLLISIONTYPE", "PEDCOUNT"], axis=1)
df2.head()
```

Out[8]:

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

| | | | | | | | | | | | |
|--|--|--|--|--|--|--|--|--|--|--|--|
| | | | | | | | | | | | |
| | | | | | | | | | | | |

In [9]:

```
#Getting the type of each column
df2.dtypes
```

Out[9]:

```
SEVERITYCODE      int64
LONGITUDE         float64
LATITUDE          float64
PERSONCOUNT      int64
VEHCOUNT          int64
JUNCTIONTYPE      object
INATTENTIONIND    object
WEATHER           object
ROADCOND          object
LIGHTCOND         object
SPEEDING          object
dtype: object
```

In [10]:

```
#Getting the shape of the data frame
df2.shape
```

Out[10]:

```
(194673, 11)
```

In [11]:

```
#Getting the name of each column
df2.columns
```

Out[11]:

```
Index(['SEVERITYCODE', 'LONGITUDE', 'LATITUDE', 'PERSONCOUNT', 'VEHCOUNT',
```

```

        'JUNCTIONTYPE', 'INATTENTIONIND', 'WEATHER', 'ROADCOND', 'LIGHTCOND'
    ,
        'SPEEDING'],
    dtype='object')

```

In [12]:

```
df2.isna().sum()
```

Out[12]:

```

SEVERITYCODE      0
LONGITUDE        5334
LATITUDE         5334
PERSONCOUNT     0
VEHCOUNT         0
JUNCTIONTYPE     6329
INATTENTIONIND   164868
WEATHER          5081
ROADCOND         5012
LIGHTCOND        5170
SPEEDING         185340
dtype: int64

```

In [13]:

```

#Returning the objects containing counts of unique values
df2['WEATHER'].value_counts()

```

Out[13]:

```

Clear              111135
Raining            33145
Overcast           27714
Unknown            15091
Snowing            907
Other               832
Fog/Smog/Smoke     569
Sleet/Hail/Freezing Rain  113
Blowing Sand/Dirt   56
Severe Crosswind    25
Partly Cloudy       5
Name: WEATHER, dtype: int64

```

Handling of missing Values

In [14]:

```

# Replacing NaN value by Unknown
df2['WEATHER'].replace(np.NaN, "Unknown", inplace=True)
df2.head()

```

Out[14]:

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

| | | | | | | | | | | | |
|--|--|--|--|--|--|--|--|--|--|--|--|
| | | | | | | | | | | | |
| | | | | | | | | | | | |

In [15]:

```
# Replacing Unknown and Other by Clear, the most frequent value of the column
encoding_WEATHER = {"WEATHER":
                    {"Clear": 1,
                     "Unknown": 1,
                     "Other": 1,
                     "Raining": 2,
                     "Overcast": 3,
                     "Snowing": 4,
                     "Fog/Smog/Smoke": 5,
                     "Sleet/Hail/Freezing Rain": 6,
                     "Blowing Sand/Dirt": 7,
```

```

        "Severe Crosswind": 8,
        "Partly Cloudy": 9}}
df2.replace(encoding_WEATHER, inplace=True)
df2['WEATHER'].value_counts()

```

Out[15]:

```

1    132139
2     33145
3     27714
4       907
5       569
6       113
7        56
8        25
9         5
Name: WEATHER, dtype: int64

```

In [16]:

```
df2['SPEEDING'].value_counts()
```

Out[16]:

```

Y     9333
Name: SPEEDING, dtype: int64

```

In [17]:

```

# Replacing NaN value by N
df2['SPEEDING'].replace(np.NaN, "N", inplace=True)

```

In [18]:

```

encoding_SPEEDING = {"SPEEDING":
                      {"Y": 1,
                       "N": 0,
                      }}
df2.replace(encoding_SPEEDING, inplace=True)
df2['SPEEDING'].value_counts()

```

Out[18]:

```

0    185340
1     9333
Name: SPEEDING, dtype: int64

```

In [19]:

```
df2['LIGHTCOND'].value_counts()
```

Out[19]:

```

Daylight          116137
Dark - Street Lights On    48507
Unknown           13473
Dusk               5902
Dawn               2502
Dark - No Street Lights    1537
Dark - Street Lights Off   1199
Other              235
Dark - Unknown Lighting    11

```

Name: LIGHTCOND, dtype: int64

In [20]:

```
# Replacing NaN value by Unknown
df2['LIGHTCOND'].replace(np.NaN, "Unknown", inplace=True)
```

In [21]:

```
# Replacing Unknown and Other by Daylight, the most frequent value of the column
```

```
encoding_LIGHTCOND = {"LIGHTCOND":
                        {"Daylight": 0,
                         "Unknown": 0,
                         "Other": 0,
                         "Dark - Street Lights On": 1,
                         "Dusk": 1,
                         "Dawn": 1,
                         "Dark - No Street Lights": 1,
                         "Dark - Street Lights Off": 1,
                         "Dark - Unknown Lighting": 1,
                        }}
df2.replace(encoding_LIGHTCOND, inplace=True)
df2['LIGHTCOND'].value_counts()
```

Out[21]:

```
0    135015
1     59658
Name: LIGHTCOND, dtype: int64
```

In [22]:

```
df2['ROADCOND'].value_counts()
```

Out[22]:

```
Dry          124510
Wet           47474
Unknown       15078
Ice            1209
Snow/Slush    1004
Other          132
Standing Water  115
Sand/Mud/Dirt   75
Oil             64
```

Name: ROADCOND, dtype: int64

In [23]:

```
# Replacing NaN value by Unknown
df2['ROADCOND'].replace(np.NaN, "Unknown", inplace=True)
```

In [24]:

```
# Replacing Unknown and Other by Dry, the most frequent value of the column
encoding_ROADCOND = {"ROADCOND":
```

```
                        {"Dry": 1,
                         "Unknown": 1,
                         "Other": 1,
```

```

        "Wet": 2,
        "Ice": 3,
        "Snow/Slush": 4,
        "Standing Water": 5,
        "Sand/Mud/Dirt": 6,
        "Oil": 7,
    }}
df2.replace(encoding_ROADCOND, inplace=True)
df2['ROADCOND'].value_counts()

```

Out [24]:

```

1    144732
2     47474
3      1209
4       1004
5        115
6         75
7         64

```

Name: ROADCOND, dtype: int64

In [25]:

```
df2['JUNCTIONTYPE'].value_counts()
```

Out [25]:

```

Mid-Block (not related to intersection)      89800
At Intersection (intersection related)        62810
Mid-Block (but intersection related)          22790
Driveway Junction                           10671
At Intersection (but not related to intersection)  2098
Ramp Junction                               166
Unknown                                      9

```

Name: JUNCTIONTYPE, dtype: int64

In [26]:

```

# Replacing NaN value by Unknown
df2['JUNCTIONTYPE'].replace(np.NaN, "Unknown", inplace=True)

```

In [27]:

```

# Replacing Unknown by Mid-Block (not related to intersection), the most fr
equent value of the column
encoding_JUNCTIONTYPE = {"JUNCTIONTYPE":
    {"Mid-Block (not related to intersection)": 1,
    "Unknown": 1,
    "At Intersection (intersection related)": 2,
    "Mid-Block (but intersection related)": 3,
    "Driveway Junction": 4,
    "At Intersection (but not related to intersect
ion)": 5,
    "Ramp Junction": 6,
    }}
df2.replace(encoding_JUNCTIONTYPE, inplace=True)

```

```
df2['JUNCTIONTYPE'].value_counts()
```

Out[27]:

```
1    96138
2    62810
3    22790
4    10671
5     2098
6      166
Name: JUNCTIONTYPE, dtype: int64
```

In [28]:

```
df2['LONGITUDE'].value_counts()
```

Out[28]:

```
-122.332653    265
-122.344896    254
-122.328079    252
-122.344997    239
-122.299160    231
...
-122.322768     1
-122.288680     1
-122.405699     1
-122.323578     1
-122.343898     1
Name: LONGITUDE, Length: 23563, dtype: int64
```

In [29]:

```
# NaN values are placed by the mean values of the column
avg_LONGITUDE = df2["LONGITUDE"].astype("float").mean(axis=0)
print("Average of LONGITUDE:", avg_LONGITUDE)
df2['LONGITUDE'].replace(np.NaN, avg_LONGITUDE, inplace=True)
Average of LONGITUDE: -122.33051843904114
```

In [30]:

```
df2['LATITUDE'].value_counts()
```

Out[30]:

```
47.708655    265
47.717173    254
47.604161    252
47.725036    239
47.579673    231
...
47.556705     1
47.709101     1
47.513899     1
47.565438     1
47.563521     1
Name: LATITUDE, Length: 23839, dtype: int64
```

In [31]:

```
# NaN values are placed by the mean values of the column
avg_LATITUDE = df2["LATITUDE"].astype("float").mean(axis=0)
print("Average of LATITUDE:", avg_LATITUDE)
df2['LATITUDE'].replace(np.NaN, avg_LATITUDE, inplace=True)

Average of LATITUDE: 47.619542517688615
```

In [32]:

```
# 1:prop damage    2:injury
df2['SEVERITYCODE'].value_counts()
```

Out[32]:

```
1    136485
2     58188
Name: SEVERITYCODE, dtype: int64
```

In [33]:

```
df2['INATTENTIONIND'].value_counts()
```

Out[33]:

```
Y    29805
Name: INATTENTIONIND, dtype: int64
```

In [34]:

```
#Replacing NaN value by N
df2['INATTENTIONIND'].replace(np.NaN, "N", inplace=True)
```

In [36]:

```
df2['INATTENTIONIND'].value_counts()
```

Out[36]:

```
N    164868
Y     29805
Name: INATTENTIONIND, dtype: int64
```

In [37]:

```
encoding_INATTENTIONIND = {"INATTENTIONIND":
                            {"Y": 1,
                             "N": 0,
                             }}
df2.replace(encoding_INATTENTIONIND, inplace=True)
```

In [38]:

```
df2["INATTENTIONIND"].value_counts()
```

Out[38]:

```
0    164868
1     29805
Name: INATTENTIONIND, dtype: int64
```

In [39]:

```
df2['SPEEDING'].value_counts()
```

Out[39]:

```
0    185340
1     9333
Name: SPEEDING, dtype: int64
```

In [40]:


```
df2.isna().sum()
```

Out[40]:

```
SEVERITYCODE      0
LONGITUDE         0
LATITUDE          0
PERSONCOUNT      0
VEHCOUNT         0
JUNCTIONTYPE      0
INATTENTIONIND    0
WEATHER           0
ROADCOND          0
LIGHTCOND         0
SPEEDING          0
dtype: int64
```

In [41]:

```
#Tabulating the first five rows
df2.head()
```

Out[41]:

[illegible]

[illegible]

| | | | | | | | | | | | |
|--|--|--|--|--|--|--|--|--|--|--|--|
| | | | | | | | | | | | |
| | | | | | | | | | | | |
| | | | | | | | | | | | |

In [42]:

```
df2.columns
```

Out[42]:

```
Index(['SEVERITYCODE', 'LONGITUDE', 'LATITUDE', 'PERSONCOUNT', 'VEHCOUNT',
      'JUNCTIONTYPE', 'INATTENTIONIND', 'WEATHER', 'ROADCOND', 'LIGHTCOND'
      ,
      'SPEEDING'],
      dtype='object')
```

In [43]:

```
df2.shape
```

Out[43]:

```
(194673, 11)
```

Methodology

After the features are selected, they are employed for an explanatory data analysis to figure out more about their effects. The focus is on identifying the feature conditions that have a bigger effect on the severity which leads to injuries. To do so, the dataset is filtered further and the corresponding values of features are sorted.

In the next step, the features are processed for predictive modeling analysis. 4 machine learning models are created using the classification techniques as listed below:

- K-Nearest Neighbors (KNN)
- Decision Tree
- Logistic Regression
- Random Forest

The created models are tested and then evaluated based on their accuracy score to find the more accurate model.

Exploratory Data Analysis

The map with markers of the accident locations in Seattle

In [44]:

```
#Installing Folium Package for mapping
#!conda install -c conda-forge folium=0.5.0 --yes
import folium
import io
from PIL import Image
```

```
print('Folium installed and imported!')
```

Folium installed and imported!

In [45]:

```
#Visualizing 300 data points on the map
limit = 300
df_m1 = df2[["LATITUDE", "LONGITUDE"]]
df_m2 = df_m1.iloc[0:limit, :]
df_m2.head()
```

Out[45]:

| | LATITUDE | LONGITUDE |
|---|-----------|-------------|
| 0 | 47.703140 | -122.323148 |
| 1 | 47.647172 | -122.347294 |
| 2 | 47.607871 | -122.334540 |

| | LATITUDE | LONGITUDE |
|----------|-----------------|------------------|
| 3 | 47.604803 | -122.334803 |
| 4 | 47.545739 | -122.306426 |

In [46]:

```
# Instantiate a feature group for the incidents in the dataframe

Seattle_map = folium.Map(location=[47.6062, -122.3321], zoom_start=12)

incidents = folium.map.FeatureGroup()

# loop through the 300 points and add each to the incidents feature group
for lat, lng, in zip(df_m2.LATITUDE, df_m2.LONGITUDE):
    incidents.add_child(
        folium.CircleMarker(
            [lat, lng],
            radius=5, # define how big you want the circle markers to be
            color='yellow',
            fill=True,
            fill_color='blue',
            fill_opacity=0.6
        )
    )

# add incidents to map
Seattle_map.add_child(incidents)
Seattle_map.save("figure3_seattle-map.html")
```

In [47]:

```
#Selecting the severity code of 2 i.e. with injuries and making another data frame for this purpose
Sev_2 = df2.loc[df2['SEVERITYCODE']==2]
Sev_2.head()
```

Out[47]:

[illegible]

[illegible]

Relationship between the weather conditions and the accident severity with injury

In [48]:

```
Sev_2_w = Sev_2['WEATHER'].value_counts()  
Sev_2_w
```

Out[48]:

```
1    37856  
2    11176  
3     8745  
5     187  
4      171  
6       28  
7       15  
8        7  
9        3  
Name: WEATHER, dtype: int64
```

In [89]:

```
labels = 'Clear', 'Raining', 'Overcast', 'Other'  
sizes = [37856, 11176, 8745, sum(Sev_2_w[3:9])]  
explode = (0.1,0, 0, 0)  
fig1, ax1 = plt.subplots(figsize=(15,7))  
ax1.pie(sizes, explode=explode, labels=labels, autopct='%1.1f%%',shadow=False, startangle=90)  
ax1.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.  
plt.title('Effect of Weather Conditions on the Severity with Injury', y=1.05)  
plt.savefig("image/figure4_weather-to-severity2.png")
```

Relationship between the person count and the accident severity with injury

In [50]:

```
Sev_2_p = Sev_2['PERSONCOUNT'].value_counts()  
Sev_2_p
```

Out[50]:

```
2    27811  
3    13461  
4     6295  
1     3296  
5     2969  
0     1762  
6     1357  
7      637  
8      284  
9      129  
10      74  
11      33
```



```

12      20
13      12
17       8
15       7
14       7
16       5
22       2
19       2
34       2
23       1
32       1
28       1
27       1
25       1
24       1
48       1
37       1
54       1
39       1
20       1
18       1
81       1
29       1
30       1
Name: PERSONCOUNT, dtype: int64

```

In [90]:

```

labels = 2, 3, 4, 1, 5, 0, 6, '>6'
sizes = [27811, 13461, 6295, 3296, 2969, 1762, 1357, sum(Sev_2_p[3:9])]
explode = (0.1, 0, 0, 0, 0, 0, 0, 0)
fig1, ax1 = plt.subplots(figsize=(15,7))
ax1.pie(sizes, explode=explode, labels=labels, autopct='%1.1f%%',shadow=False, startangle=90)
ax1.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.
plt.title('Effect of Person count on the Severity with Injury', y=1)
plt.savefig("image/figure5_personcount-to-severity2.png")

```

Relationship between the vehicle count and the accident severity with injury

In [52]:

```

Sev_2_v = Sev_2['VEHCOUNT'].value_counts()
Sev_2_v

```

Out[52]:

```

2      35949
1      14105
3       5470
0       1227
4       1078

```

```

5      261
6      60
7      22
9       6
8       5
11      3
10      2
Name: VEHCOUNT, dtype: int64

```

In [91]:

```

labels = 2, 1, 3, 0, 4, '>4'
sizes = [35949, 14105, 5470, 1227, 1078, sum(Sev_2_p[5:12])]
explode = (0.1, 0, 0, 0, 0, 0)
fig1, ax1 = plt.subplots(figsize=(15,7))
ax1.pie(sizes, explode=explode, labels=labels, autopct='%1.1f%%', shadow=False, startangle=90)
ax1.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.
plt.title('Effect of Vehicle count on the Severity with Injury', y=1.05)
plt.savefig("image/figure6_vehiclecount-to-severity2.png")

```

Relationship between the junction type and the accident severity with injury

In [54]:

```

Sev_2_j = Sev_2['JUNCTIONTYPE'].value_counts()
Sev_2_j

```

Out[54]:

```

2      27174
1      19806
3       7297
4       3234
5        623
6         54
Name: JUNCTIONTYPE, dtype: int64

```

In [92]:

```

labels = 'At Intersection_intersection related', 'Mid-Block_not intersection related', 'Mid-Block with intersection', 'Driveway Junction', 'Other'
sizes = [27174, 19806, 7297, 3234, sum(Sev_2_j[4:6])]
explode = (0.1, 0, 0, 0, 0)
fig1, ax1 = plt.subplots(figsize=(15,7))
ax1.pie(sizes, explode=explode, labels=labels, autopct='%1.1f%%', shadow=False, startangle=90)
ax1.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.
plt.title('Effect of Junction type on the Severity with Injury', y=1.05)
plt.savefig("image/figure7_junction-to-severity2.png")

```

Relationship between the inattention and the accident severity with injury

In [56]:

```
Sev_2_i = Sev_2['INATTENTIONIND'].value_counts()
Sev_2_i
```

Out[56]:

```
0    47791
1    10397
Name: INATTENTIONIND, dtype: int64
```

In [95]:

```
labels = 'No', 'Yes'
sizes = [47791, 10397]
explode = (0.1, 0)
fig1, ax1 = plt.subplots(figsize=(15,7))
ax1.pie(sizes, explode=explode, labels=labels, autopct='%1.1f%%', shadow=False, startangle=90)
ax1.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.
plt.title('Effect of Inattention on the Severity with Injury', y=1)
plt.savefig("image/figure8_inattention-to-severity2.png")
```

Relationship between the road conditions and the accident severity with injury

In [58]:

```
Sev_2_r = Sev_2['ROADCOND'].value_counts()
Sev_2_r
```

Out[58]:

```
1    41916
2    15755
3      273
4      167
5       30
7       24
6       23
Name: ROADCOND, dtype: int64
```

In [96]:

```
labels = 'dry', 'wet', 'ic, sand, oil, standing water'
sizes = [41916, 15755, sum(Sev_2_r[2:7])]
explode = (0.1, 0, 0)
fig1, ax1 = plt.subplots(figsize=(15,7))
ax1.pie(sizes, explode=explode, labels=labels, autopct='%1.1f%%', shadow=False, startangle=90)
ax1.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.
plt.title('Effect of Road condition on the Severity with Injury', y=1.05)
plt.savefig("image/figure9_roadcondition-to-severity2.png")
```

Relationship between the light conditions and the accident severity with injury

In [60]:

```
Sev_2_1 = Sev_2['LIGHTCOND'].value_counts()  
Sev_2_1
```

Out[60]:

```
0    40291  
1    17897  
Name: LIGHTCOND, dtype: int64
```

In [97]:

```
labels = 'day light', 'dark'  
sizes = [40291, 17897]  
explode = (0.1, 0)  
fig1, ax1 = plt.subplots(figsize=(15,7))  
ax1.pie(sizes, explode=explode, labels=labels, autopct='%1.1f%%', shadow=False, startangle=90)  
ax1.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.  
plt.title('Effect of Light condition on the Severity with Injury', y=1)  
plt.savefig("image/figure10_lightcondition-to-severity2.png")
```

Relationship between the speeding and the accident severity with injury

In [62]:

```
Sev_2_s = Sev_2['SPEEDING'].value_counts()  
Sev_2_s
```

Out[62]:

```
0    54657  
1     3531  
Name: SPEEDING, dtype: int64
```

In [98]:

```
labels = 'No', 'Yes'  
sizes = [54657, 3531]  
explode = (0.1, 0)  
fig1, ax1 = plt.subplots(figsize=(15,7))  
ax1.pie(sizes, explode=explode, labels=labels, autopct='%1.1f%%', shadow=False, startangle=90)  
ax1.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.  
plt.title('Effect of Speeding condition on the Severity with Injury', y=1)  
plt.savefig("image/figure11_speeding-to-severity2.png")
```

Modeling, Testing and Evaluation

In [64]:

```
#Making a new data frame  
Feature = df2[['LONGITUDE', 'LATITUDE', 'PERSONCOUNT', 'VEHCOUNT',  
              'JUNCTIONTYPE', 'INATTENTIONIND', 'WEATHER', 'ROADCOND', 'LIGHTCOND',  
              'SPEEDING']]
```

In [66]:

```
X = Feature
print(X)
```

| | LONGITUDE | LATITUDE | PERSONCOUNT | VEHCOUNT | JUNCTIONTYPE | \ |
|--------|-------------|-----------|-------------|----------|--------------|-----|
| 0 | -122.323148 | 47.703140 | 2 | 2 | 2 | |
| 1 | -122.347294 | 47.647172 | 2 | 2 | 1 | |
| 2 | -122.334540 | 47.607871 | 4 | 3 | 1 | |
| 3 | -122.334803 | 47.604803 | 3 | 3 | 1 | |
| 4 | -122.306426 | 47.545739 | 2 | 2 | 2 | |
| ... | ... | ... | ... | ... | ... | ... |
| 194668 | -122.290826 | 47.565408 | 3 | 2 | 1 | |
| 194669 | -122.344526 | 47.690924 | 2 | 2 | 1 | |
| 194670 | -122.306689 | 47.683047 | 3 | 2 | 2 | |
| 194671 | -122.355317 | 47.678734 | 2 | 1 | 2 | |
| 194672 | -122.289360 | 47.611017 | 2 | 2 | 1 | |

| | INATTENTIONIND | WEATHER | ROADCOND | LIGHTCOND | SPEEDING |
|--------|----------------|---------|----------|-----------|----------|
| 0 | 0 | 3 | 2 | 0 | 0 |
| 1 | 0 | 2 | 2 | 1 | 0 |
| 2 | 0 | 3 | 1 | 0 | 0 |
| 3 | 0 | 1 | 1 | 0 | 0 |
| 4 | 0 | 2 | 2 | 0 | 0 |
| ... | ... | ... | ... | ... | ... |
| 194668 | 0 | 1 | 1 | 0 | 0 |
| 194669 | 1 | 2 | 2 | 0 | 0 |
| 194670 | 0 | 1 | 1 | 0 | 0 |
| 194671 | 0 | 1 | 1 | 1 | 0 |
| 194672 | 0 | 1 | 2 | 0 | 0 |

[194673 rows x 10 columns]

In [67]:

```
y = df2['SEVERITYCODE'].values
print(y)
```

[2 1 1 ... 2 2 1]

Normalize Data

In [68]:

```
X= preprocessing.StandardScaler().fit(X).transform(X)
X[0:5]
```

Out[68]:

```
array([[ 0.24930404,  1.50948129, -0.33020207,  0.12553783,  0.24566547,
        -0.42518348,  1.91514317,  1.3926872 , -0.66472702, -0.22440165],
       [-0.56747188,  0.49889979, -0.33020207,  0.12553783, -0.81596734,
        -0.42518348,  0.64986567,  1.3926872 ,  1.50437693, -0.22440165],
       [-0.1360361 , -0.21073866,  1.15576451,  1.7102107 , -0.81596734,
        -0.42518348,  1.91514317, -0.53629605, -0.66472702, -0.22440165],
       [-0.14494267, -0.26614566,  0.41278122,  1.7102107 , -0.81596734,
```

```

-0.42518348, -0.61541182, -0.53629605, -0.66472702, -0.22440165],
[ 0.81495737, -1.33262277, -0.33020207,  0.12553783,  0.24566547,
-0.42518348,  0.64986567,  1.3926872 , -0.66472702, -0.22440165]])

```

In [69]:

```

#Splitting the data into train-test sets
X_train, X_test, y_train, y_test = train_test_split( X, y, test_size=0.2, r
andom_state=4)
print ('Train set:', X_train.shape,  y_train.shape)
print ('Test set:', X_test.shape,  y_test.shape)

Train set: (155738, 10) (155738,)
Test set: (38935, 10) (38935,)

```

K-Nearest Neighbours (KNN)

In []:

```

#Finding the best k
Ks = 10
mean_acc = np.zeros((Ks-1))
std_acc = np.zeros((Ks-1))
#ConfustionMx = [];
for n in range(1,Ks):
    #Train Model and Predict
    knNeigh = KNeighborsClassifier(n_neighbors = n).fit(X_train,y_train)
    yhat = knNeigh.predict(X_test)
    mean_acc[n-1] = metrics.accuracy_score(y_test, yhat)
    std_acc[n-1] = np.std(yhat==y_test)/np.sqrt(yhat.shape[0])

```

In [84]:

```

#Plot model accuracy for Different number of Neighbors
plt.plot(range(1,Ks),mean_acc,'g')
plt.fill_between(range(1,Ks),mean_acc - 1 * std_acc,mean_acc + 1 * std_acc,
alpha=0.10)
plt.legend(('Accuracy ', '+/- 3xstd'))
plt.ylabel('Accuracy ')
plt.xlabel('Number of Nabors (K)')
plt.tight_layout()
plt.savefig("image/figure12_KNN.png")

```

In [85]:

```

print( "Best k =", mean_acc.argmax()+1)

Best k = 8

```

In [70]:

```

#Building the model using the best k
k=8
knNeigh= KNeighborsClassifier(n_neighbors = k).fit(X_train,y_train)
knNeigh

```

Out[70]:

```

KNeighborsClassifier(n_neighbors=8)

```

In [71]:

```
#Evalaution
yhat = knNeigh.predict(X_test)
KNN_accuracy_score = metrics.accuracy_score(y_test, yhat)
print("K-Nearest Neighbours Accuray: ", KNN_accuracy_score)

K-Nearest Neighbours Accuray: 0.7329395145755747
```

Decision Tree

In [72]:

```
#Modeling
DTree = DecisionTreeClassifier(criterion="entropy", max_depth = 4)
DTree.fit(X_train,y_train)
```

Out[72]:

```
DecisionTreeClassifier(criterion='entropy', max_depth=4)
```

In [73]:

```
#Prediction
yhat = DTree.predict(X_test)
yhat
```

Out[73]:

```
array([1, 1, 1, ..., 1, 1, 1])
```

In [74]:

```
#Evalaution
DT_accuracy_score = metrics.accuracy_score(y_test, yhat)
print("DecisionTrees's Accuracy: ", DT_accuracy_score)

DecisionTrees's Accuracy: 0.7429562090663927
```

Logistic Regression

In [75]:

```
#Modeling
LR = LogisticRegression(C=0.01, solver='liblinear').fit(X_train,y_train)
LR
```

Out[75]:

```
LogisticRegression(C=0.01, solver='liblinear')
```

In [76]:

```
#Prediction
yhat = LR.predict(X_test)
yhat
```

Out[76]:

```
array([1, 1, 1, ..., 1, 1, 1])
```

In [77]:

```
#Evalaution
LR_accuracy_score = metrics.accuracy_score(y_test, yhat)
print("Logistic Regression's Accuracy: ", LR_accuracy_score)

Logistic Regression's Accuracy: 0.7030949017593425
```

Random Forest

In [78]:

```
#Modeling
clf=RandomForestClassifier(n_estimators=100)
```

```
clf.fit(X_train,y_train)
```

Out[78]:

```
RandomForestClassifier()
```

In [79]:

```
#Prediction
```

```
yhat =clf.predict(X_test)
```

In [80]:

```
#Evalaution
```

```
RF_accuracy_score = metrics.accuracy_score(y_test, yhat)
```

```
print("Random Forest's Accuracy: ", RF_accuracy_score)
```

```
Random Forest's Accuracy: 0.7114164633363298
```

Visualize important features using Random Forest

In [99]:

```
# Create a list of feature names
```

```
feat_labels = ['SEVERITYCODE', 'LONGITUDE', 'LATITUDE','PERSONCOUNT',  
               'VEHCOUNT', 'JUNCTIONTYPE', 'INATTENTIONIND',  
               'WEATHER', 'ROADCOND', 'LIGHTCOND', 'SPEEDING']
```

```
# Set the target for the prediction
```

```
target='SEVERITYCODE'
```

```
# Create arrays for the features and the response variable
```

```
# set X and y
```

```
y = df2[target]
```

```
X = df2.drop(target, axis=1)
```

```
# Split the data set into training and testing data sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra  
ndom_state=21, stratify=y)
```

```
# Create a random forest classifier
```

```
clf = RandomForestClassifier(n_estimators=100)
```

```
# Train the classifier
```

```
clf.fit(X_train, y_train)
```

```
feature_imp = pd.Series(clf.feature_importances_,index=X.columns).sort_valu  
es(ascending=False)
```

```
# Creating a bar plot, displaying only the top k features
```

```
k=10
```

```
sns.barplot(x=feature_imp[:10], y=feature_imp.index[:k])
```

```
# Add labels to your graph
```

```
plt.xlabel('Feature Importance Score')
```

```
plt.ylabel('Features')
```

```
plt.title("Visualizing Important Features")
```



```
plt.savefig("image/figure13_Random-Forest.png")
```

Results and Discussion

Plot the accuracy score versus algorithm

In [102]:

```
algo_lst = ['K-Nearest Neighbors', 'Decision Trees', 'Logistic Regression', 'Random Forest']
```

```
accuracy_lst = [KNN_accuracy_score, DT_accuracy_score, LR_accuracy_score, RF_accuracy_score]
```

```
# Generate a list of ticks for y-axis
```

```
y_ticks=np.arange(len(algo_lst))
```

```
#Combine the list of algorithms and list of accuracy scores into a dataframe, sort the value based on accuracy score
```

```
df_acc=pd.DataFrame(list(zip(algo_lst, accuracy_lst)), columns=['Algorithm', 'Accuracy_Score']).sort_values(by=['Accuracy_Score'],ascending = True)
```

```
# Make a plot
```

```
ax=df_acc.plot.barh('Algorithm', 'Accuracy_Score', align='center',legend=False)
```

```
# Add the data label on to the plot
```

```
for i in ax.patches:
```

```
    # get_width pulls left or right; get_y pushes up or down
```

```
    ax.text(i.get_width()+0.1, i.get_y()+0.2, str(round(i.get_width(),2)),
fontsize=10)
```

```
# Set the limit, labels, ticks and title
```

```
plt.xlim(0,1.1)
```

```
plt.xlabel('Accuracy Score')
```

```
plt.yticks(y_ticks, df_acc['Algorithm'], rotation=0)
```

```
plt.title('Accuracy Score versus Algorithm')
```

```
plt.savefig("image/figure14_algorithm-score.png")
```

Comparing the score of accuracies obtained by the algorithms K-Nearest Neighbors, Decision Tree, Logistic Regression, and Random Forest, the decision tree has been proved to give better accuracy.

During the modeling with the K-Nearest Neighbors classifier, it was observed that the computer required much more time. But it took less time to execute the decision tree modeling. This can also represent better effectiveness and compatibility of the decision tree for handling this given dataset.

Conclusion and Outlook

In this study, supervised machine learning is employed to predict car accident severity. The imbalanced dataset is firstly balanced, and the raw data is understood and prepared in different steps to be used for the predictive modeling analysis. In parallel, an explanatory data analysis is done to gain more insight into the relationship between the features and the severity of the accidents.

Four machine learning algorithms (K-Nearest Neighbors, Decision Trees, Logistic Regression, and Random Forest) are applied in which the decision tree has shown better compatibility with the dataset, resulting in higher accuracy (0.74).

One idea for future work can be developing the decision tree machine learning model to improve its accuracy further. Adding more data to the dataset can help to compensate for the missing values. Gathering more data about other parameters such as the age of the drivers can also help to gain a more detailed insight into the car accident severity.